

# What is Kubernetes?

- Started by Google in 2014. First released in 2015
- Kubernetes is an open-source system for automating deployment, scaling, and management of containerized applications.
- Donated to Cloud Native Computing Foundation
- 100% open source
- Written in Go language

# Certified Kubernetes Administrator Program

- The Certified Kubernetes Administrator (CKA) program was created by the Cloud Native Computing Foundation (CNCF), in collaboration with The Linux Foundation, to help develop the Kubernetes ecosystem.
- The purpose of the CKA program is to provide assurance that CKAs have the skills, knowledge, and competency to perform the responsibilities of Kubernetes administrators
- Performance-based problems to be solved in a command line.
- Time 3 Hours
- Cost \$300 with 1 free retake
- Questions 24
- 74% passing

# CKA domains & weights

- Installation, Configuration & Validation 12%
- Application Lifecycle Management 8%
- Core Concepts 19%
- Networking 11%
- Scheduling 5%
- Security 12%
- Cluster Maintenance 11%
- Logging / Monitoring 5%
- Storage 7%
- Troubleshooting 10%

# CKA Exam Curriculum

## Installation, Configuration & Validation – 12%

- Design a Kubernetes Cluster
- Install Kubernetes Masters and Nodes
- Configure a highly-available Kubernetes cluster
- Know where to get the Kubernetes release binaries
- Provision underlying infrastructure to deploy a Kubernetes cluster
- Choose a network solution
- Choose your Kubernetes infrastructure configuration
- Analyze end-to-end test results
- Run Node end-to-end Tests
- Install and use kubeadm to install, configure, and manage clusters

# CKA Exam Curriculum

## Core Concepts – 19%

- Understand the Kubernetes API primitives
- Understand the Kubernetes cluster architecture
- Understand Services and other network primitives

## Application Lifecycle Management – 8%

- Understand deployments and how to perform rolling update and rollbacks
- Know various ways to configure applications
- Know how to scale applications
- Understand the primitives necessary to create a self-healing application

# CKA Exam Curriculum

## Networking – 11%

- Understand the networking configuration on the cluster nodes
- Understand Pod networking concepts
- Understand Service Networking
- Deploy and configure network load balancer
- Know how to use Ingress rules
- Know how to configure and use the cluster DNS
- Understand CNI

# CKA Exam Curriculum

## Scheduling – 5%

- Use label selectors to schedule Pods
- Understand the role of DaemonSets
- Understand how resource limits can affect Pod scheduling
- Understand how to run multiple schedulers and how to configure Pods to use them
- Manually schedule a pod without a scheduler
- Display scheduler events

# CKA Exam Curriculum

## Security – 12%

- Know how to configure authentication and authorization
- Understand Kubernetes security primitives
- Know how to configure network policies
- Create and manage TLS certificates for cluster components
- Work with images securely
- Define security contexts
- Secure persistent key value store



# CKA Exam Curriculum

## Cluster Maintenance – 11%

- Understand Kubernetes cluster upgrade process
- Facilitate operating system upgrades
- Implement backup and restore methodologies

## Logging / Monitoring – 5%

- Understand how to monitor all cluster components
- Understand how to monitor applications
- Manage cluster component logs
- Manage application logs

# CKA Exam Curriculum

## Storage – 7%

- Understand persistent volumes and know how to create them
- Understand access modes for volumes
- Understand persistent volume claims primitive
- Understand Kubernetes storage objects
- Know how to configure applications with persistent storage

## Troubleshooting – 10%

- Troubleshoot application failure
- Troubleshoot control plane failure
- Troubleshoot worker node failure
- Troubleshoot networking

# Certification Details

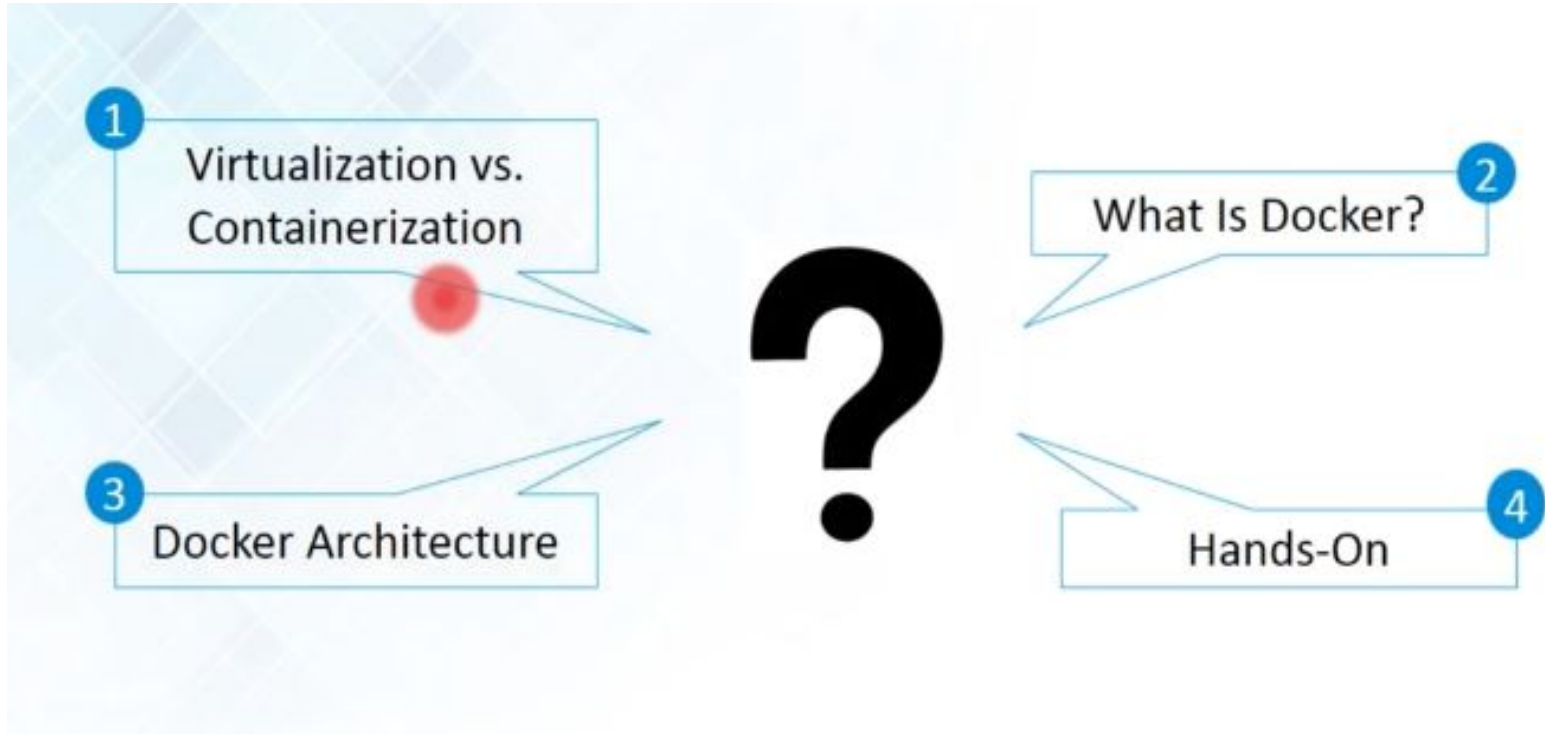
Certified Kubernetes Administrator: <https://www.cncf.io/certification/cka/>

Exam Curriculum (Topics): <https://github.com/cncf/curriculum>


Candidate Handbook: <https://www.cncf.io/certification/candidate-handbook>

Exam Tips: <http://training.linuxfoundation.org/go/Important-Tips-CKA-CKAD>

# Docker



# Docker Internals



# Docker Internals

- Containers are a method of operating system virtualization that allow you to run application and its dependencies in a resource isolated process.
- Containers allow you to easily package an application code, configurations and dependencies into easy to use building blocks and that application can be deployed quickly and consistently regardless of deployment environment

# Docker Internals

Linux kernel features that create the walls between container and other processes running on the host.

To understand Containers We have to start with **Linux Cgroup and Namespace & Union File System**.

**Linux Namespace** - Wrap a set of system resources and present them to a process to make it look like they are dedicated to that process.

# Docker Internals

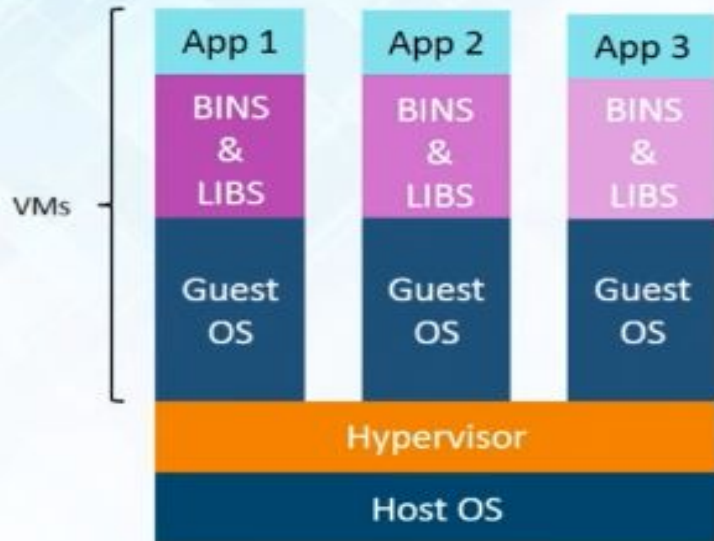
**Linux Cgroup** - Governs the isolation and usage of system resources such as cpu and memory for group of process. E.g. - If you have a application that takes up lot of cpu cycles and memory such as scientific computing application you can put the application in a cgroup to limit a CPU and memory usage.

- **Namespaces** deal with resource isolation for single process
- **Cgroup** manages resources for group of processes



# Virtualization

## Virtualization Technique



### Advantages

- Multiple OS In Same Machine
- Easy Maintenance & Recovery
- Lower Total Cost Of Ownership

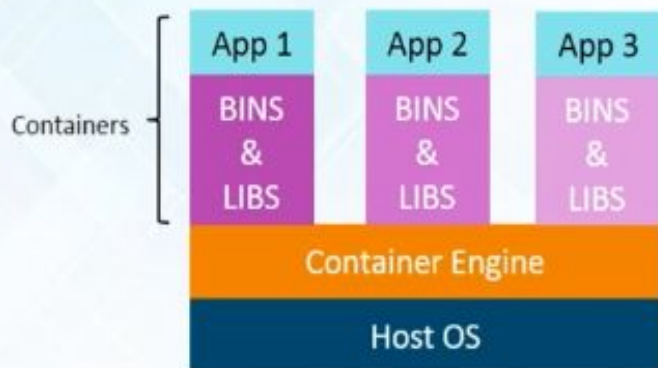
### Disadvantages

- Multiple VMs Lead To Unstable Performance
- Hypervisors Are Not As Efficient As Host OS
- Long Boot-Up Process ( [Approx. 1 Minute](#) )

# Containerization

Note: Containerization Is Just Virtualization At The OS Level

## Containerization Technique




## Advantages Over Virtualization

- Containers On Same OS Kernel Are Lighter & Smaller
- Better Resource Utilization Compared To VMs
- Short Boot-Up Process (  $1/20^{\text{th}}$  of a second )

# Docker Info

Docker is a Containerization platform which packages your **application and all its dependencies** together in the form of **Containers** so as to ensure that your application works seamlessly in any environment be it **Development or Test or Production**.



# Docker Info



**VS**



**SIZE**



**STARTUP**




**INTEGRATION**





# Docker Info

## In case of Virtual Machines



7 Gb of Memory is blocked and cannot be allotted to a new VM

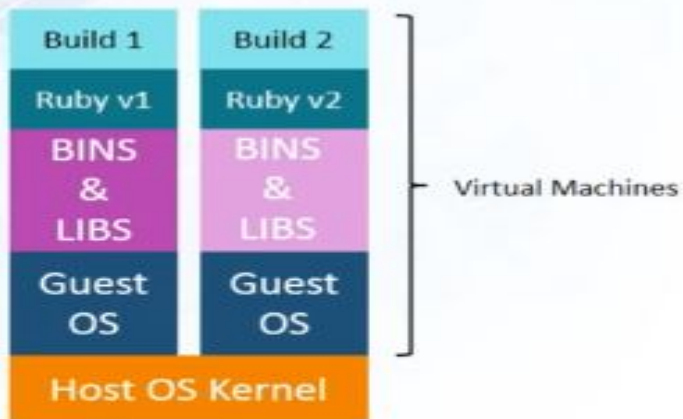
## In case of Docker



Only 9 GB memory utilized;  
7 GB can be allotted to a new Container


# Docker Info

## In case of Virtual Machines



New Builds → Multiple OS → Separate Libraries  
→ Heavy → More Time

## In case of Docker

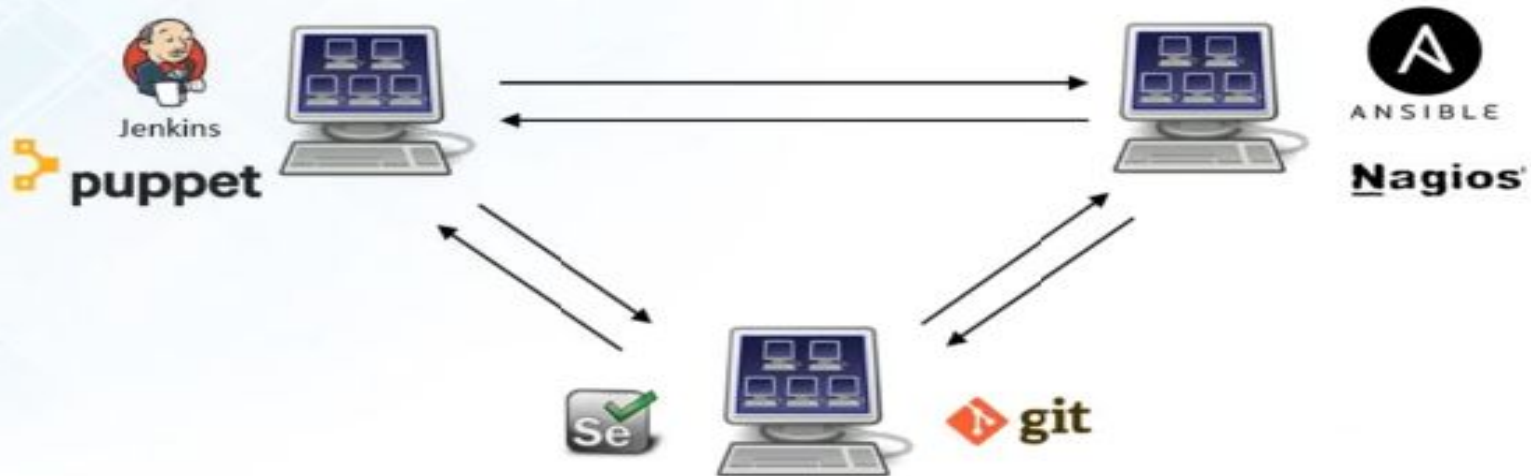


New Builds → Same OS → Separate Libraries  
→ Lightweight → Less Time

# Integration in VMs

Integration In Virtual Machines Is Possible, But:

- **Costly** Due To Infrastructure Requirements
- Not **Easily Scalable**





# Who Can Use Docker

Docker is designed to benefit both **Developers** and **System Administrators**, making it a part of many DevOps toolchains.

Developers can write code without worrying about the testing / production environment



Dev



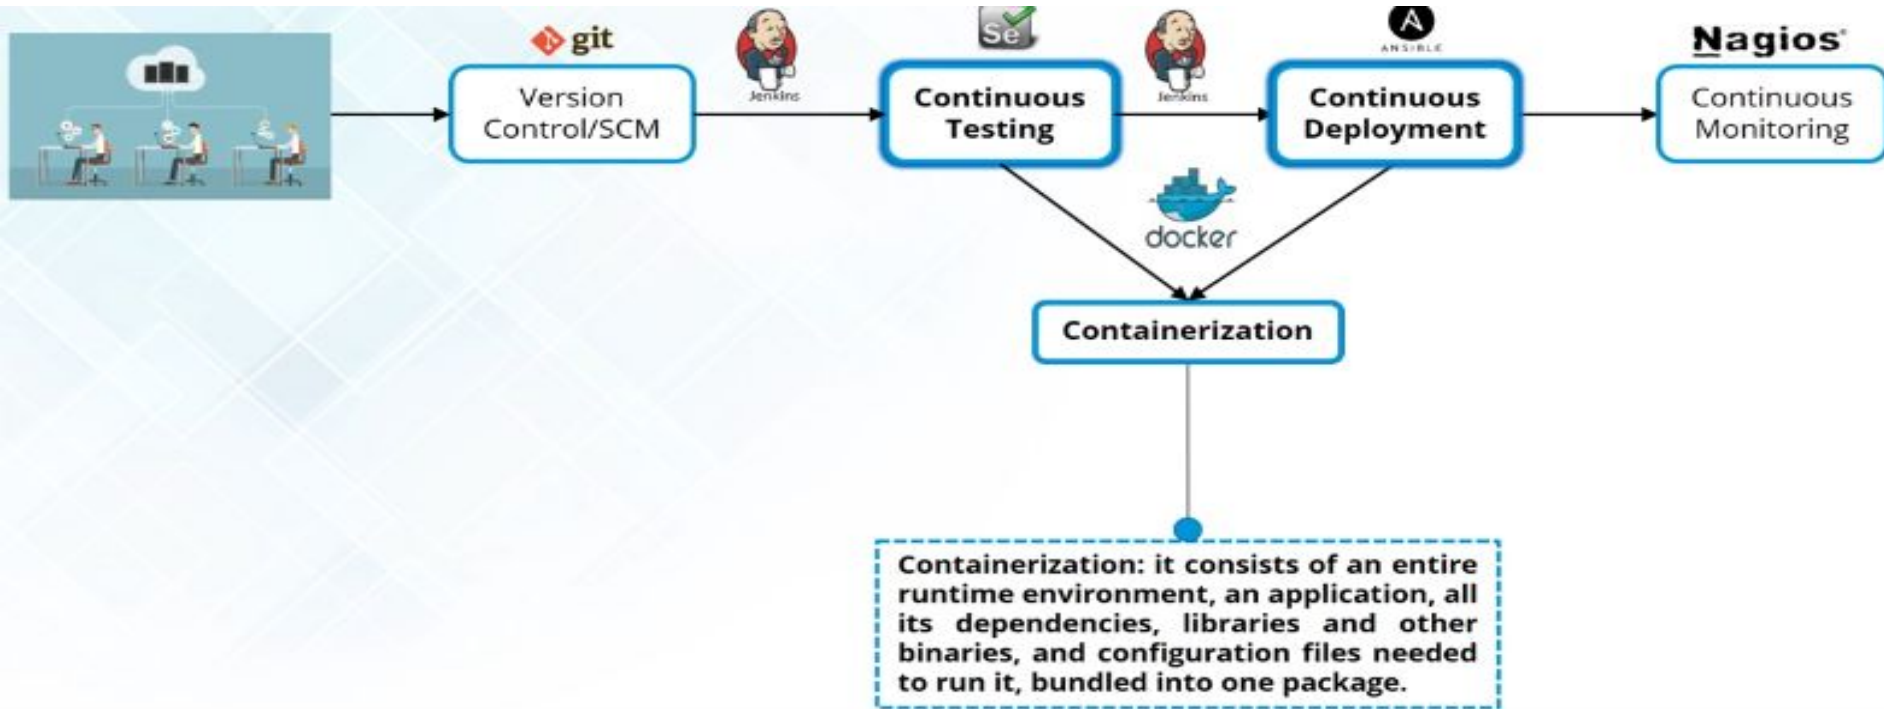
OPS

SysAdmins need not worry about Infrastructure as Docker can easily scale up / scale down the no. of systems





# How is Docker Used In Devops



# Docker Images & Containers



Docker Images

run



Docker Containers

- Read Only Template Used To Create Containers
- Built By Docker Users
- Stored In Docker Hub Or Your Local Registry

- Isolated Application Platform
- Contains Everything Needed To Run The Application
- Built From One Or More Images

# Docker Registry

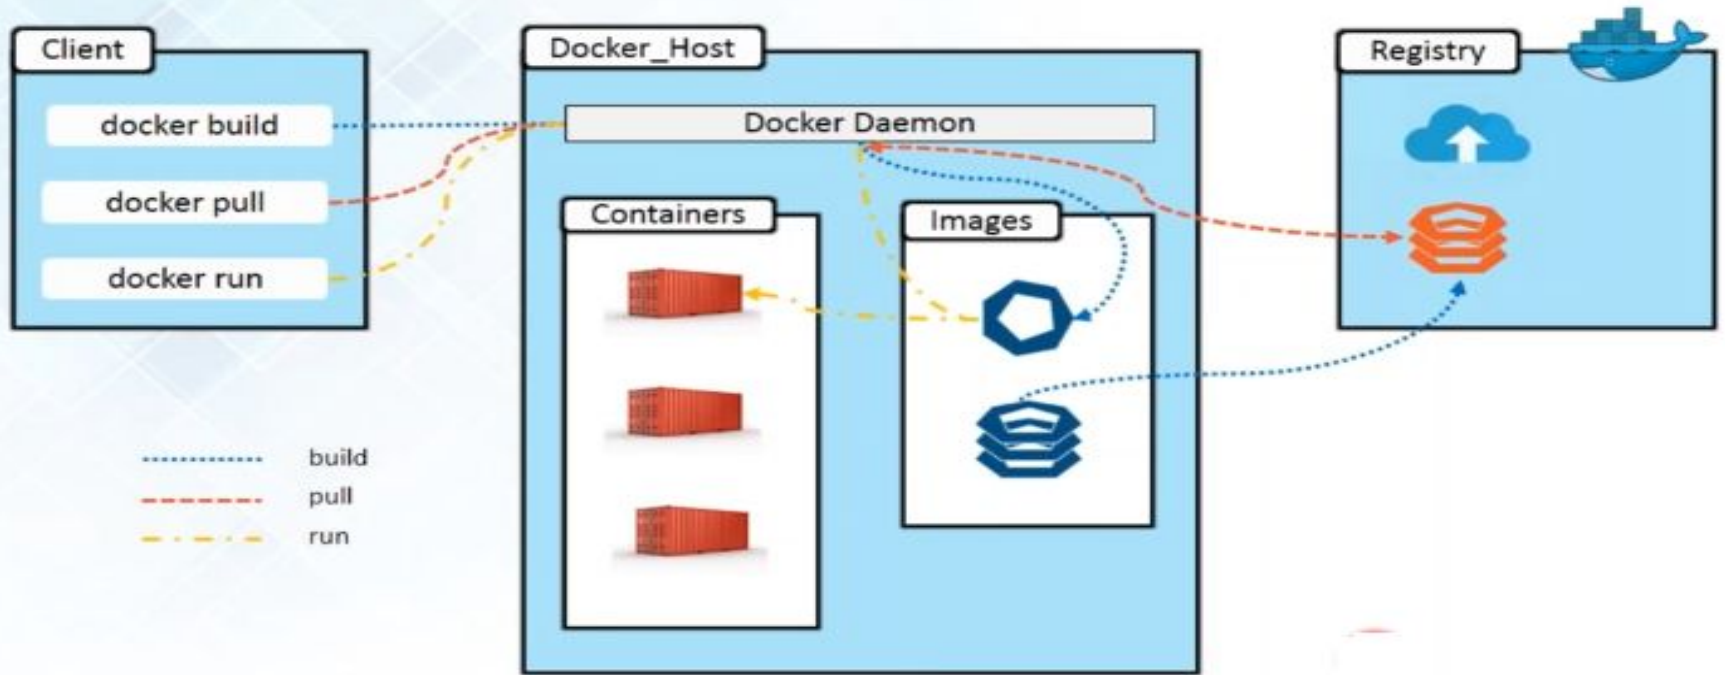
- Docker Registry is a storage component for Docker Images
- We can store the Images in either Public / Private repositories
- [Docker Hub](#) is Docker's very own cloud repository



## Why Use Docker Registries?

- Control where your images are being stored
- Integrate image storage with your in-house development workflow

# Docker Architecture



# Docker Installation

```
sudo apt-get install apt-transport-https ca-certificates curl  
software-properties-common  
  
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -  
  
sudo add-apt-repository "deb [arch=amd64]  
https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"  
  
sudo apt-get update  
  
sudo apt-get install docker-ce  
  
docker --version  
  
sudo docker run hello-world
```

# Docker Compose Installation

```
sudo apt-get -y install python-pip
```

```
sudo pip install docker-compose
```

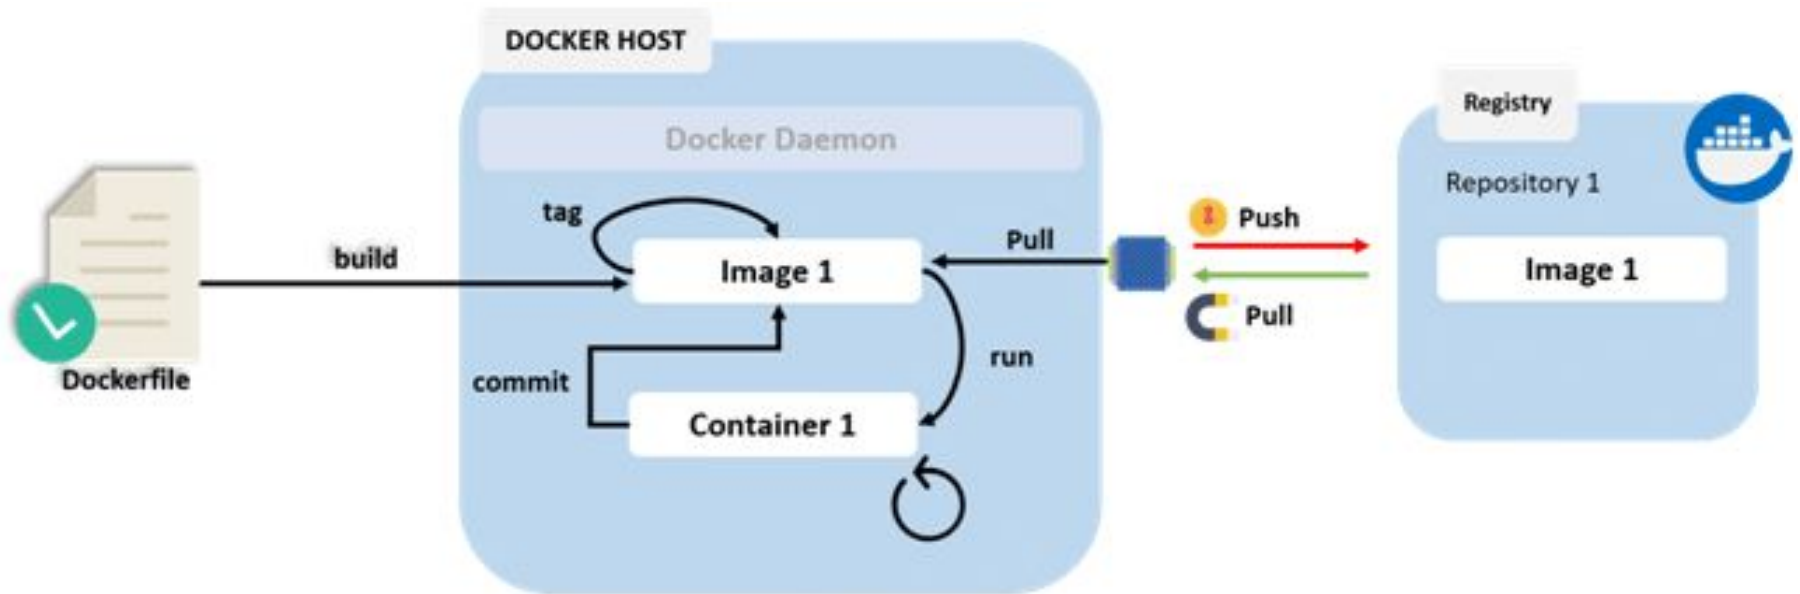
```
COMPOSE FILE: docker-compose.yml
```

**My-test:**

```
image: hello-world
```

```
docker-compose up
```

# Typical Workflow



# Basic Docker Commands

To Pull a Docker image from the Docker hub, we can use the command:

```
$ docker pull <image-name:tag>
```

To Run that image, we can use the command:

```
$ docker run <image-name:tag> or $ docker run <image-id>
```

To list down all the images in our system, we can give the command:

```
$ docker images
```

To list down all the running containers, we can use the command:

```
$ docker ps
```

To list down all the containers ( even if they are not running ), we can use the command:

```
$ docker ps -a
```



# Building Images

- Images are comprised of multiple layers
- Each layer in an Image is an Image of its own
- They comprise of a Base Image layer which is read-only
- Any changes made to an Image are saved as layers on top of the Base Image layer
- Containers are generated by running the Image layers which are stacked one above the other

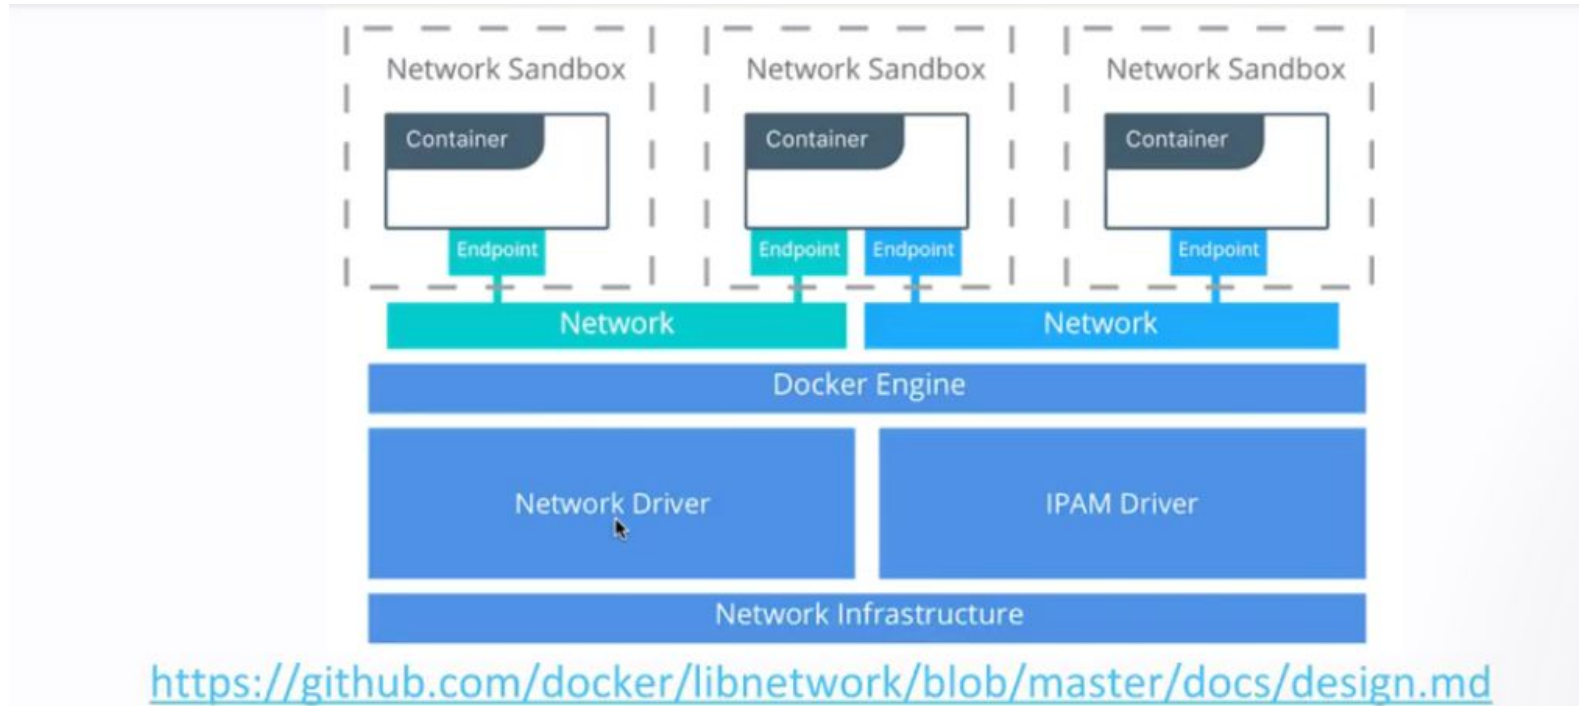


# Docker File Instructions

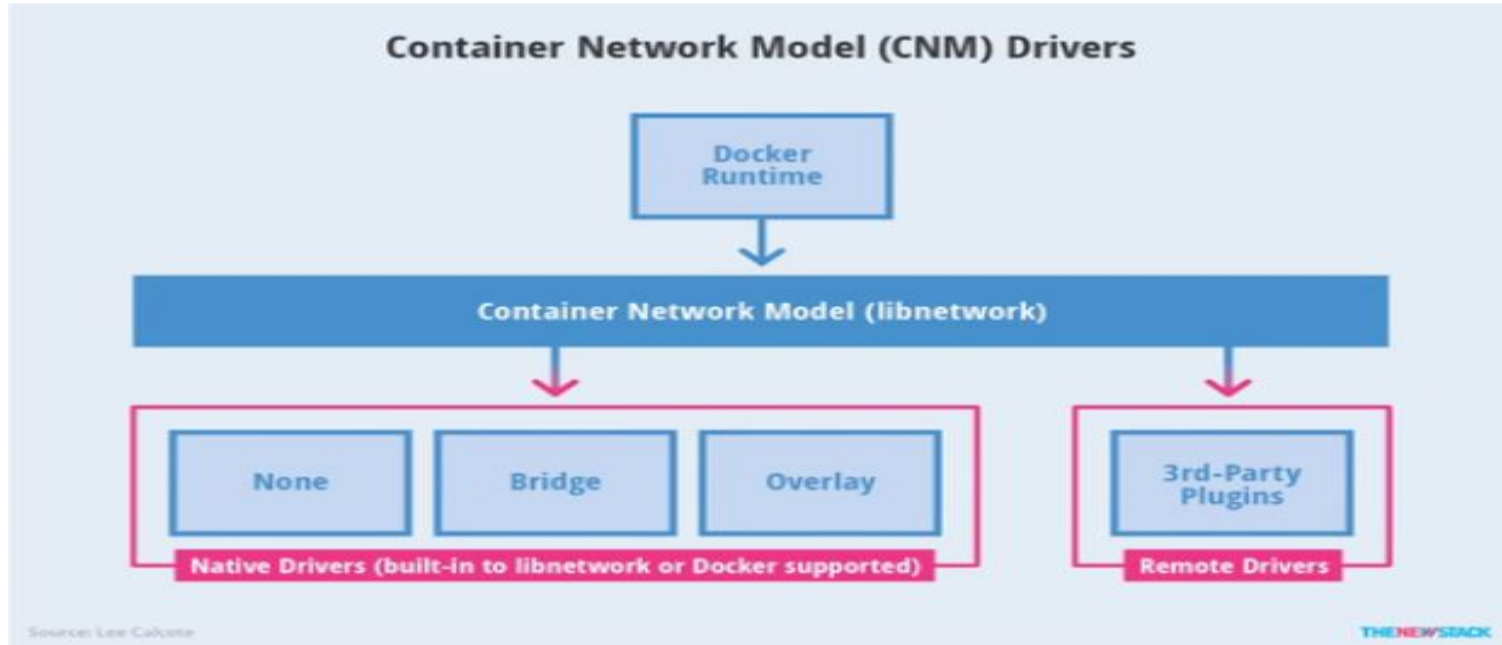
<https://docs.docker.com/engine/reference/builder/#environment-replacement>

Instruction	Comments
FROM <image>:<tag>	Tells Docker what will be the base for the image being created
ADD	Copies the files from the source on the host into the Docker image's own filesystem at the desired destination. Command is not only about copying files from the local filesystem--you can use it to get the file from the network.
COPY	It will copy new files or directories from <source path> and adds them to the filesystem of the container at the path <destination path>.
CMD / ENTRYPOINT	CMD ["executable","parameter1","parameter2"] his is the so-called exec form CMD command parameter1 parameter2 This a shell form of the instruction
RUN	Instruction will execute any commands in a new layer on top of the current image and then commit the results
EXPOSE	Instruction informs Docker that the container listens on the specified network ports at runtime.

# Container Networking Model (CNM)




# Container Networking Model

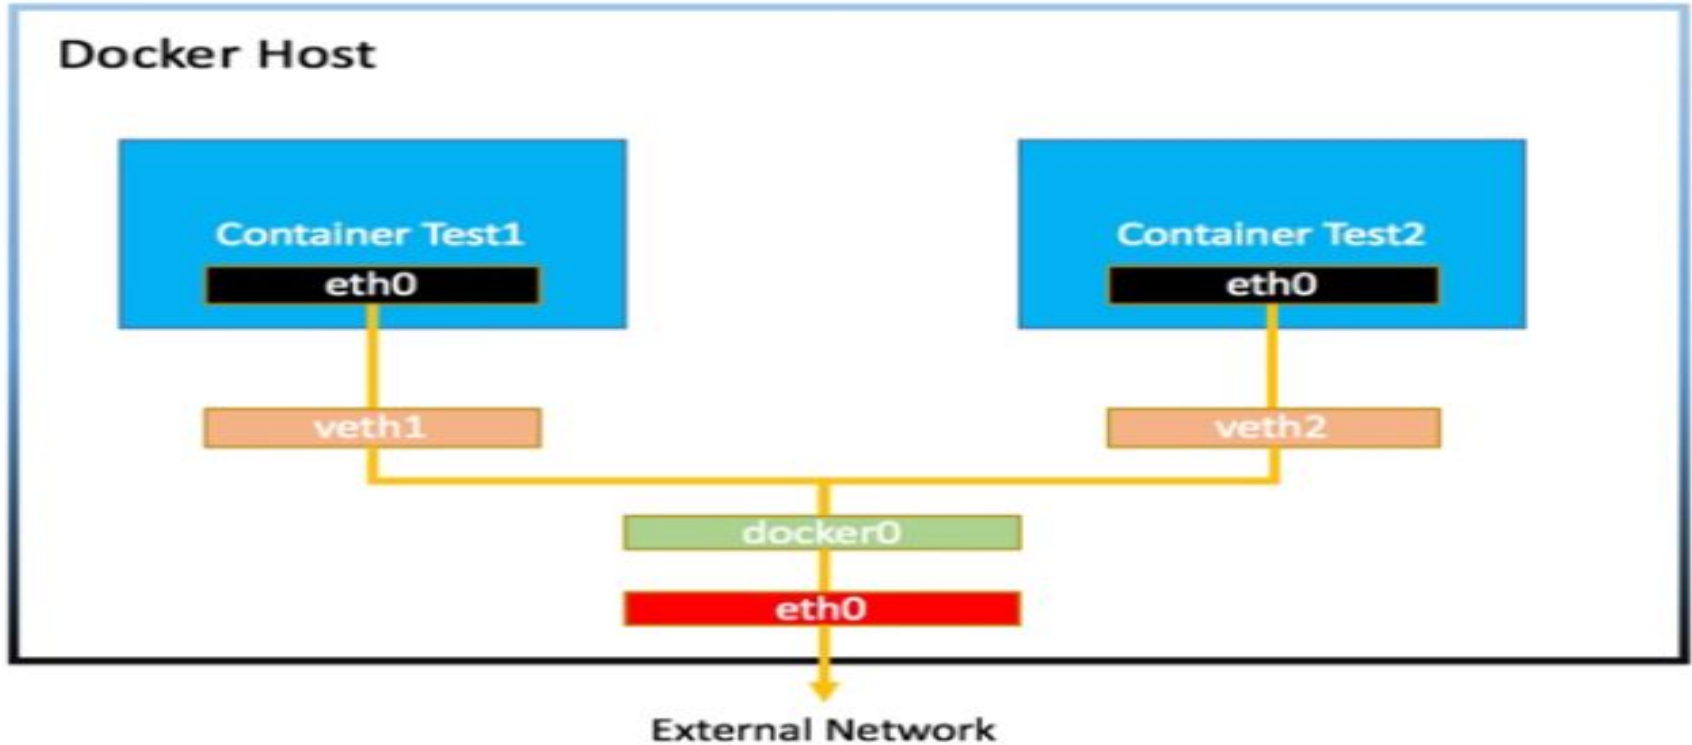


# Single Host Container Networking

- How two different containers in the same host communicate with each other?
- How the container communicate with the outside of Linux host (Internet)?
- How access container from the local docker host and outside of the docker host?

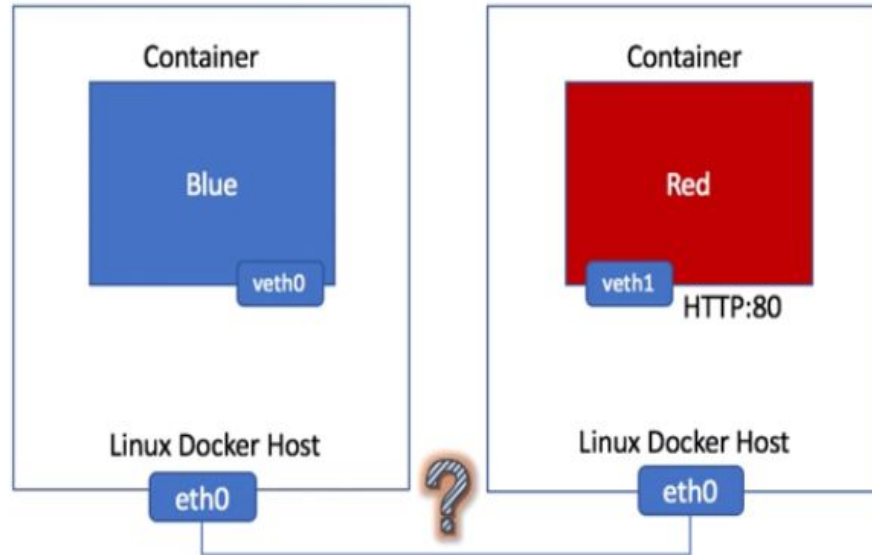


# Solution: Linux Bridge and Iptables



# Multi Host Container Networking

→ How two containers located on different docker host communicate with each other?



# Solution: Tunnel and Routing

## → Tunnel


- ◆ Docker build-in overlay network: VXLAN
- ◆ OVS: VXLAN or GRE
- ◆ Flannel: VXLAN or UDP
- ◆ Weave: VXLAN or UDP

## → Routing

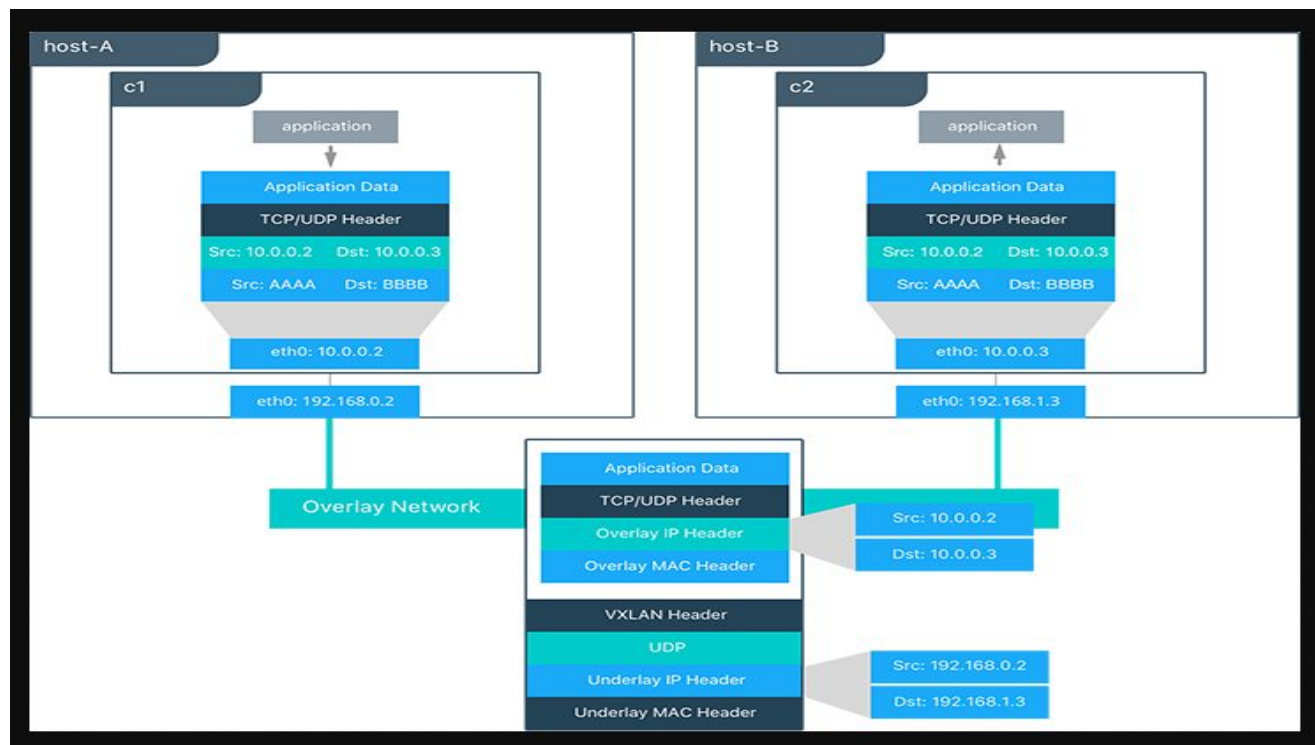
- ◆ Calico: Layer 3 routing based on BGP
- ◆ Contiv: Layer 3 routing based on BGP



# Multi Host Networking



# Multi Host Networking




# Docker Compose

# Multi-container apps are a hassle

- Build images from Dockerfiles
- Pull images from the Hub or a private registry Configure and create containers
- Start and stop containers


# Docker Compose



**Text file**



`$ docker up`



# Overview

- Compose is a tool for defining and running multi-container Docker applications. With Compose, you use a Compose file to configure your application's services. Then, using a single command, you create and start all the services from your configuration.

# app.py

```
from flask import Flask
from redis import Redis
import os
import socket

app = Flask(__name__)
redis = Redis(host=os.environ.get('REDIS_HOST', 'redis'), port=6379)

@app.route('/')
def hello():
    redis.incr('hits')
    return 'Hello Container World! I have been seen %s times and my hostname is %s.\n' % (redis.get('hits'), socket.gethostname())

if __name__ == "__main__":
    app.run(host="0.0.0.0", port=5000, debug=True)
```

# requirements.txt

flask

redis



# Dockerfile

FROM python:2.7

MAINTAINER R K "rajendrait99@gmail.com"

COPY . /app

WORKDIR /app

RUN pip install -r requirements.txt

EXPOSE 5000

CMD [ "python", "app.py" ]


# docker-compose.yml

```
version: "2"

services:
  web:
    build: .
    ports:
      - "80:5000"
    links:
      - redis
    networks:
      - compose-demo-bridge

  redis:
    image: redis
    ports: ["6379"]
    networks:
      - compose-demo-bridge

networks:
  compose-demo-bridge:
```



# Why we need Container Orchestration?

- Do you need scaling beyond one host?
- Do you need high availability?
- Are your containers truly stateless?

...

# Production requirement

- **Availability** – it just has to be working all the time, with as little downtimes as possible
- **Performance** – our server needs to handle traffic, so performance is important
- **Easy deployment & rollback**
- **Gathering logs & metrics**

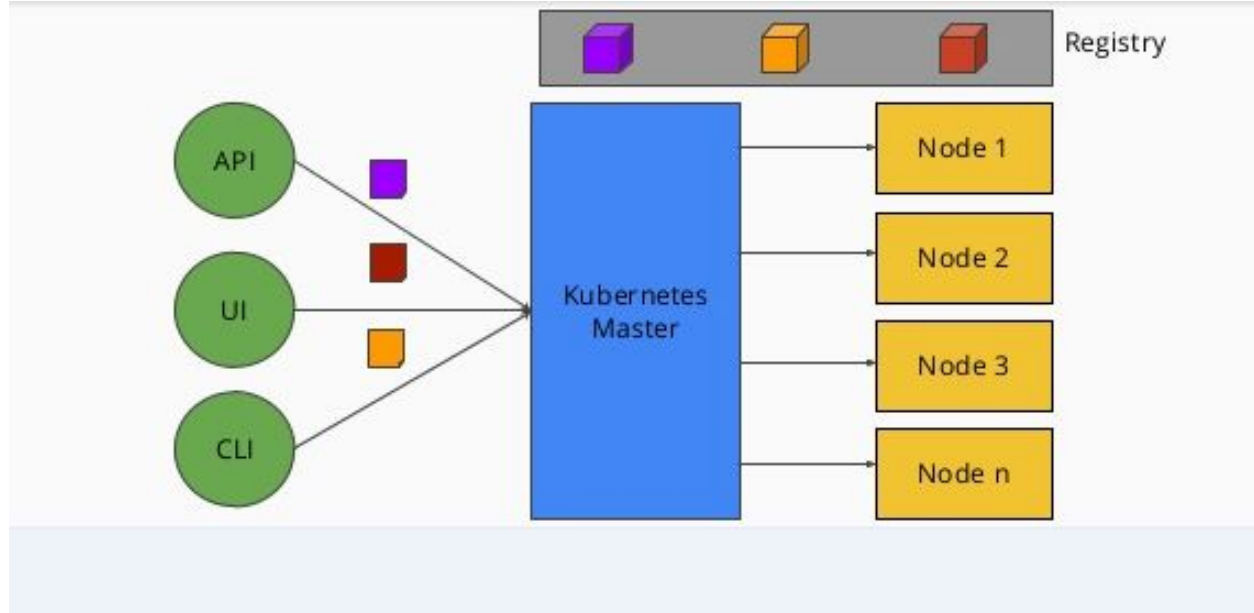
# Solutions

- **Kubernetes**
- **Docker Swarm**
- **Amazon EKS**
- **AKS**
- **GKE**

# What is Kubernetes?

- Started by Google in 2014. First released in 2015
- Kubernetes is an open-source system for automating deployment, scaling, and management of containerized applications.
- Donated to Cloud Native Computing Foundation
- 100% open source
- Written in Go language

# Kubernetes - Architecture




# Production requirement

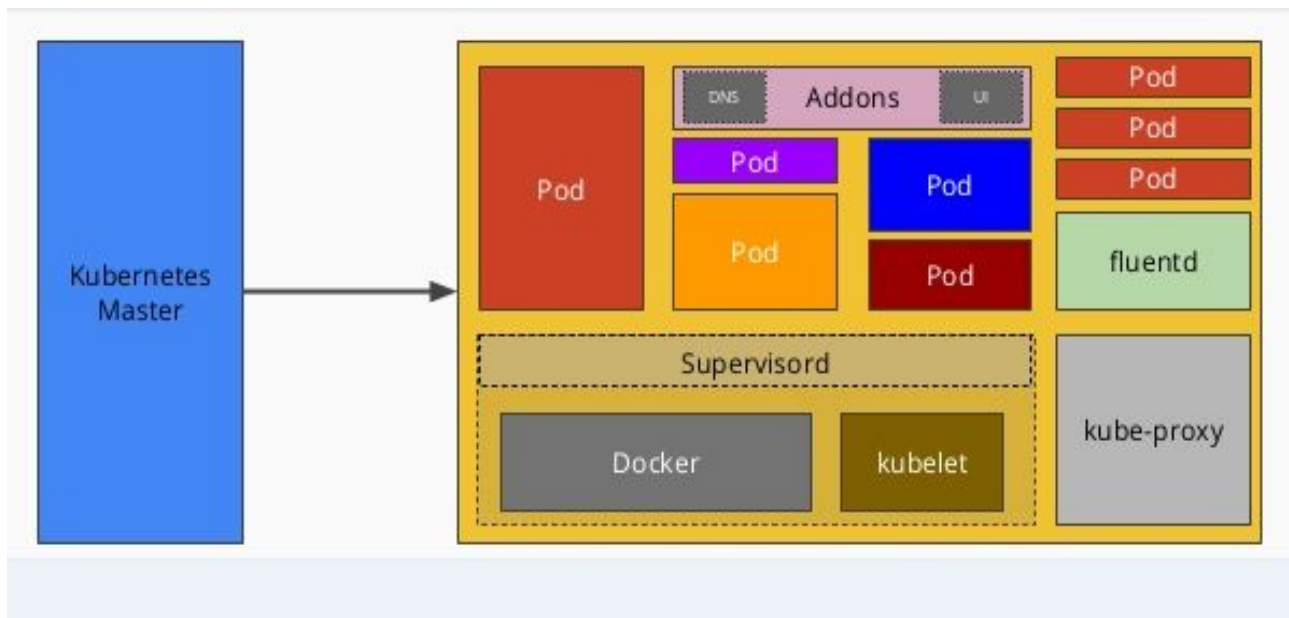
- **Scheduling:** Where do these containers run?
- **Management:** Who manages their life cycle?
- **Service Discovery:** How do they find each other?
- **Load Balancing:** How to route requests?



# Kubernetes - Architecture - Master



# Kubernetes - Architecture - Node



# Multi Node Cluster Setup

1. Run below steps on master and slave machine

- `apt-get update && apt-get install -y apt-transport-https`
- `curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -`
- `cat <<EOF > /etc/apt/sources.list.d/kubernetes.list`  
`deb http://apt.kubernetes.io/ kubernetes-xenial main`  
`EOF`

2. `apt-get update`

3. `apt-get install -y docker.io`

4. `apt-get install -y kubelet kubeadm kubectl kubernetes-cni`

# Multi Node Cluster Setup

## Init Master node and Slave Nodes

- Run below command on master node.

kubeadm init --apiserver-advertise-address=<Master IP> --pod-network-cidr=10.244.0.0/16 --skip-preflight-checks

- Copy the join command that you will get after executing above init command.  
Example: kubeadm join --token 8ccfe0.7cd3dd5f91ef9796 <Master IP>:6443
- Execute join command on all the slave nodes
- Start using cluster from Master node

**\*\*Note\*\*** : We can do this from any node where kubectl installed

- Run below command to copy configuration to default directory

```
sudo cp /etc/kubernetes/admin.conf $HOME/
```

```
sudo chown $(id -u):$(id -g) $HOME/admin.conf
```

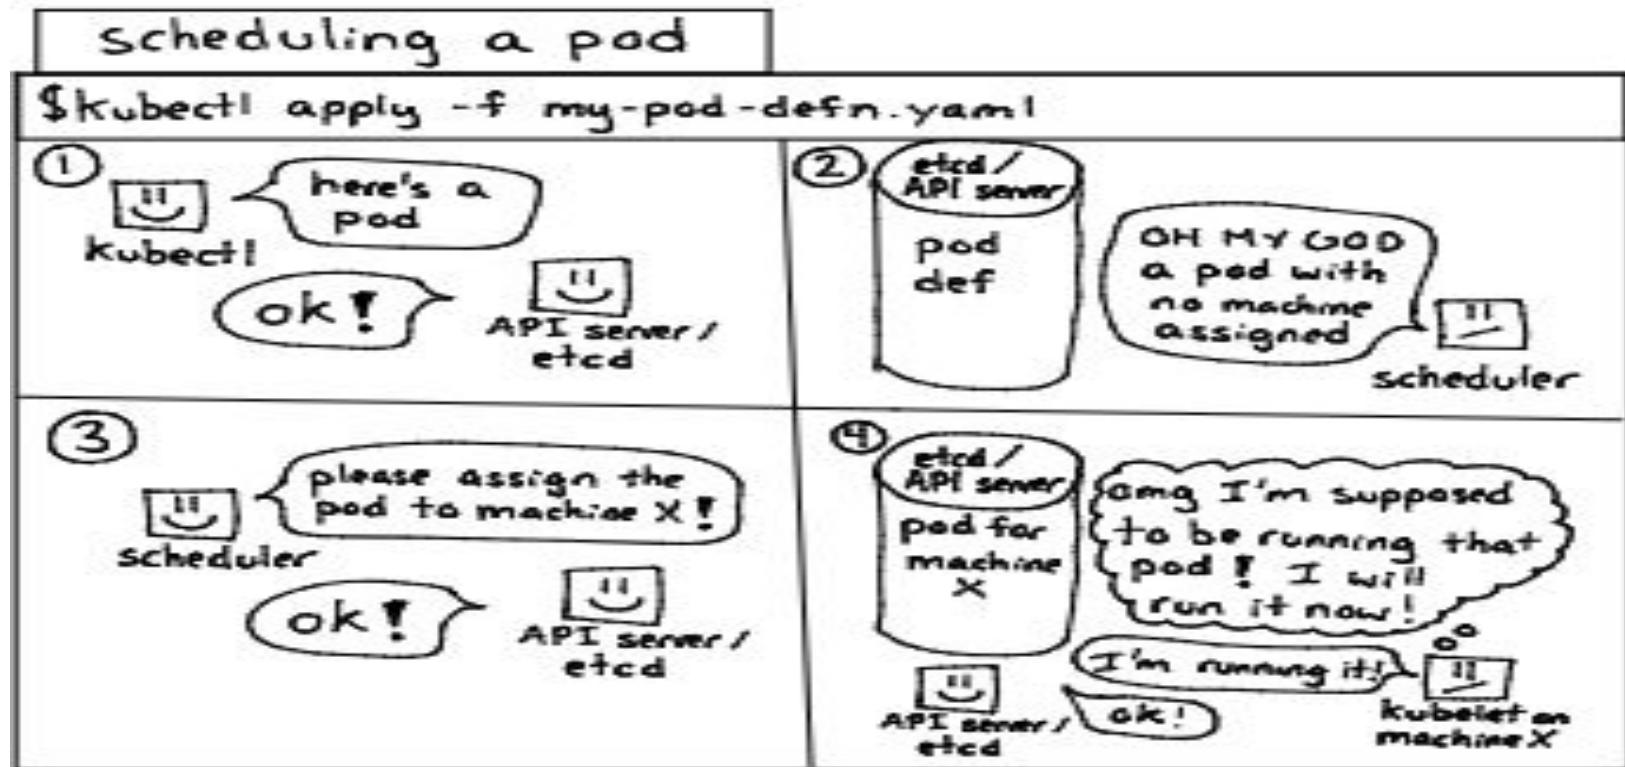
```
export KUBECONFIG=$HOME/admin.conf
```

- Setup Kubernetes Network. Run below command on master node

```
kubectl apply -f
```

<http://docs.projectcalico.org/v2.2/getting-started/kubernetes/installation/hosted/kubeadm/1.6/calico.yaml>

# Scheduling a Pod



# KIND

- **KIND** is a tool for running local Kubernetes clusters using Docker container “nodes”.
- It was primarily designed for testing Kubernetes itself, but may be used for local development or CI.
- Kubernetes **IN**side **D**ocker

# Kind Installation

```
curl -Lo ./kind https://kind.sigs.k8s.io/dl/v0.8.1/kind-linux-amd64
```

```
chmod +x ./kind
```

```
sudo mv ./kind /usr/local/bin/kind
```

# Kubectl Installation

```
curl -LO  
"https://storage.googleapis.com/kubernetes-release/release/$(curl -s  
https://storage.googleapis.com/kubernetes-release/release/stable.txt  
) /bin/linux/amd64/kubectl"
```

```
chmod +x ./kubectl
```

```
sudo mv ./kubectl /usr/local/bin/kubectl
```



# Kind config

```
# three node (two workers) cluster config
```

```
kind: Cluster
```

```
apiVersion: kind.x-k8s.io/v1alpha4
```

```
nodes:
```

- role: control-plane
- role: worker
- role: worker

# Writing YAML

```
apiVersion: v1
kind: Pod
metadata:
  name: rss-site
  labels:
    app: web
spec:
  containers:
    - name: front-end
      image: nginx
      ports:
        - containerPort: 80
    - name: rss-reader
      image: nickchase/rss-php-nginx:v1
      ports:
        - containerPort: 88
```

# Kubernetes Namespace

- Namespace as a virtual cluster inside your Kubernetes cluster
- Default namespace
- Creating Namespaces
- Viewing Namespaces
- Creating Resources in the Namespace
- Viewing resources in the Namespace
- Managing test, staging, and production environments within the same cluster

# Kubernetes - Services

- Exposing PODs to the outside world. Services are required for discovery.
- Labels and selectors are used to route to the appropriate POD
- Service Types
  - ClusterIP - Services makes internal pod accessible
  - NodePort - Allow external user traffic and load balancing
  - LoadBalancer - Service does load balancing with external load balancer

# Kubernetes - Services

```
kind: Service
apiVersion: v1
```

## metadata:

```
name: hostname-service
```

Make the service available to network requests from external clients

## spec:

```
type: NodePort
```

## selector:

```
app: echo-hostname
```

Forward requests to pods with label of this value

## ports:

```
- nodePort: 30163
```

```
port: 8080
```

```
targetPort: 80
```

### nodePort

access service via this external port number

### port

port number exposed internally in cluster


### targetPort

port that containers are listening on

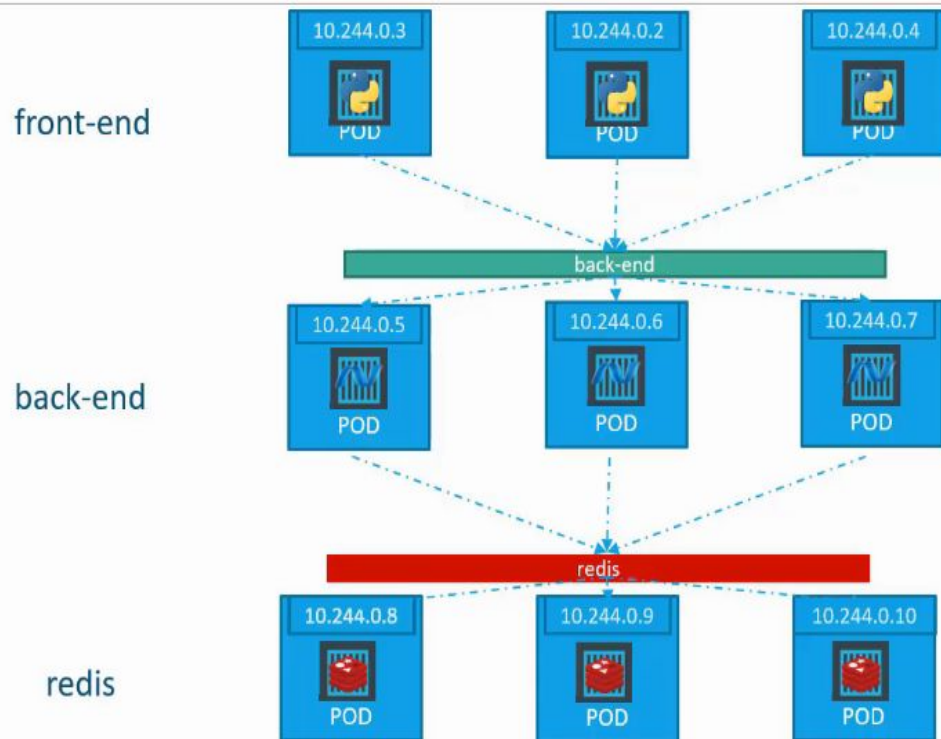
# Service - ClusterIP

- A ClusterIP service is the default Kubernetes service. It gives you a service inside your cluster that other apps inside your cluster can access. There is no external access.

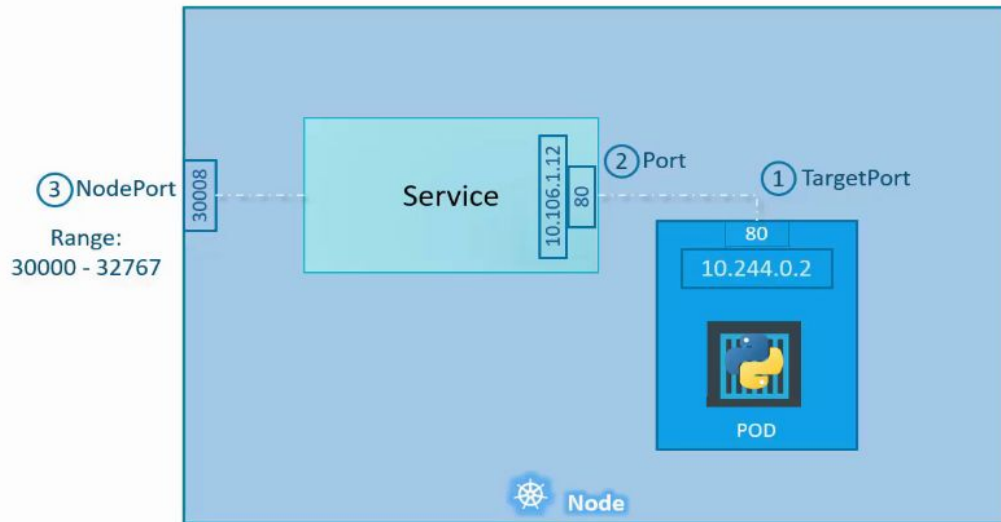
```
apiVersion: v1
kind: Service
metadata:
  name: my-internal-service
spec:
  selector:
    app: my-app
  type: ClusterIP
  ports:
    - name: http
      port: 80
      targetPort: 80
```



# ClusterIP



# Service - NodePort



service-definition.yml

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      *port: 80
      nodePort: 30008
```



# Service - NodePort

---

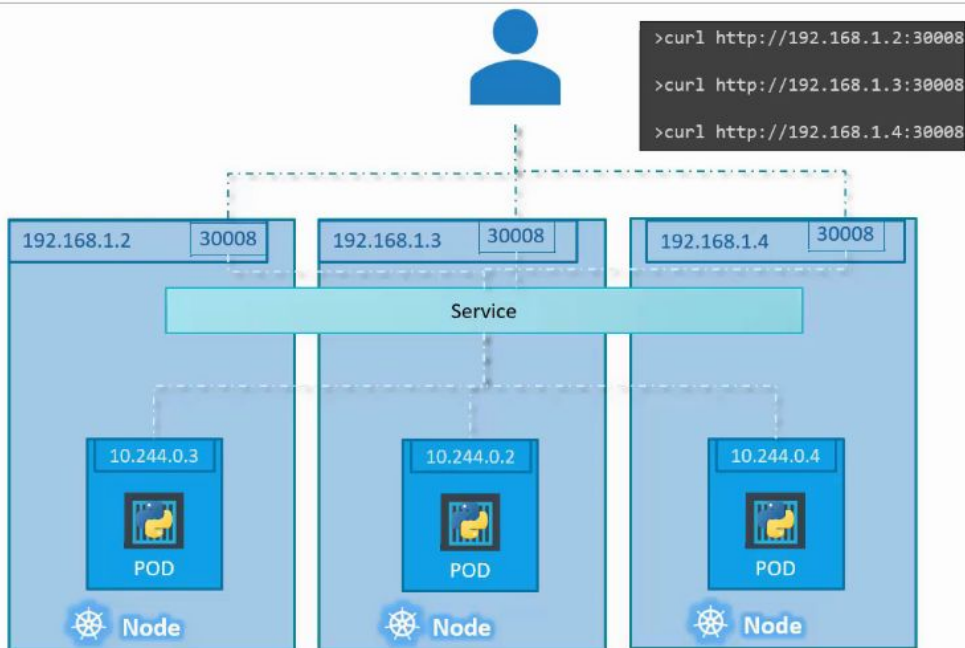
service-definition.yml

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30008
  selector:
```

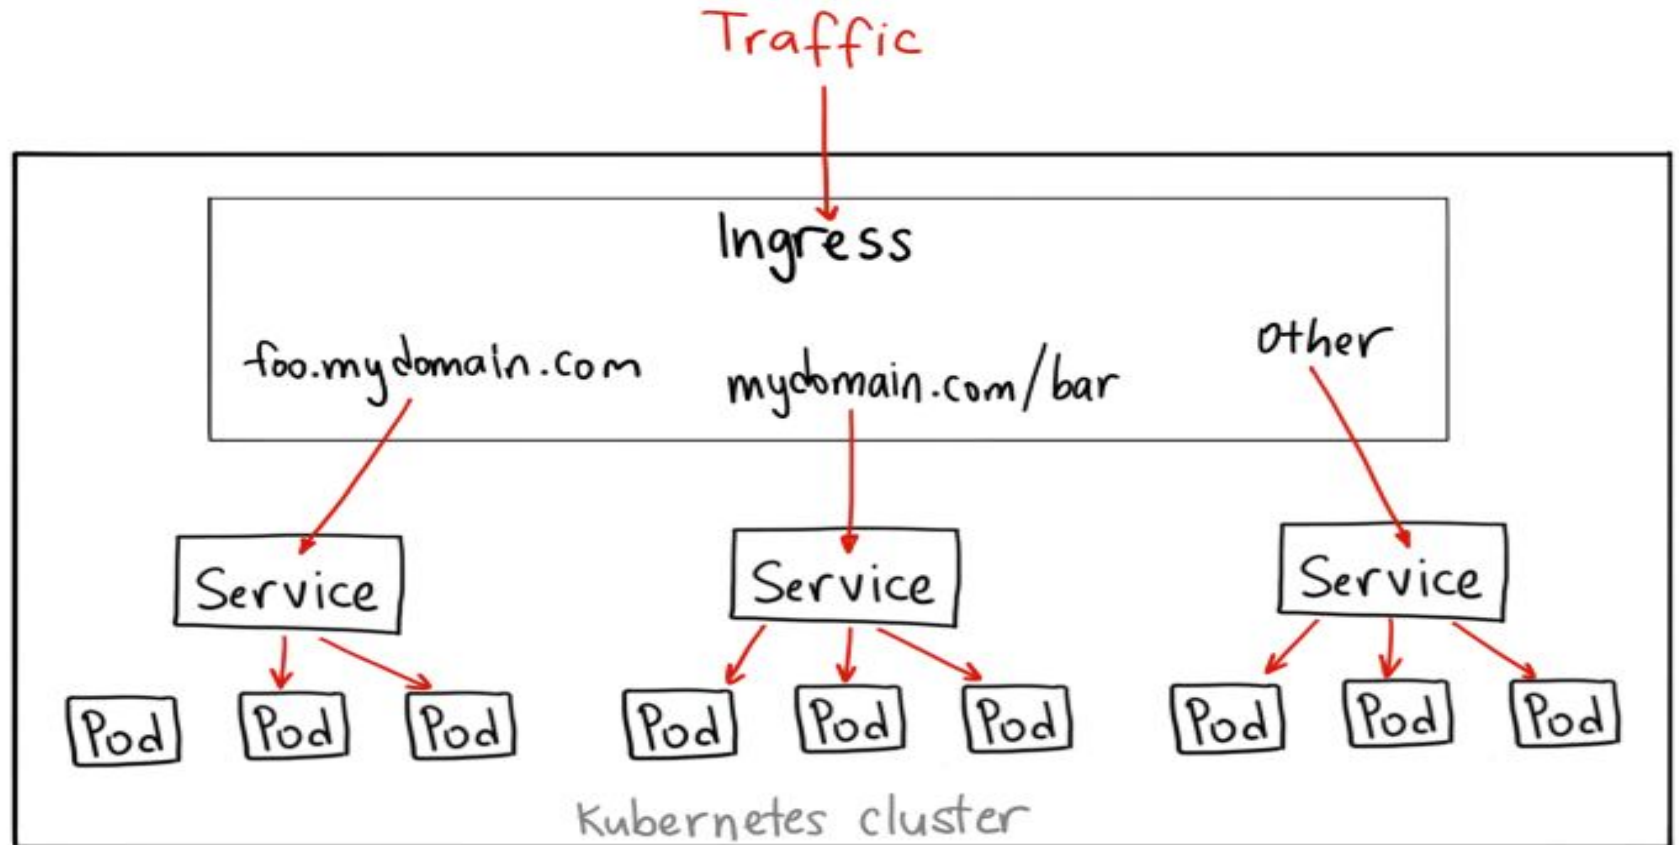
pod-definition.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end
spec:
  containers:
    - name: nginx-container
      image: nginx
```

# Service - NodePort




# Service - Ingress



# Service - Load Balancer

A LoadBalancer service is the standard way to expose a service to the internet

The big downside is that each service you expose with a LoadBalancer will get its own IP address, and you have to pay for a LoadBalancer per exposed service, which can get expensive!



# Kubernetes - Deployment

- Deployments represent a set of multiple, identical Pods with no unique identities.
- A Deployment runs multiple replicas of your application and automatically replaces any instances that fail or become unresponsive

The Kubernetes deployment object lets you:

- Deploy a replica set or pod
- Update pods and replica sets
- Rollback to previous deployment versions
- Scale a deployment
- Pause or continue a deployment

# Service - Ingress Yaml

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: my-ingress
spec:
  backend:
    serviceName: other
    servicePort: 8080
  rules:
    - host: foo.mydomain.com
      http:
        paths:
          - backend:
              serviceName: foo
              servicePort: 8080
    - host: mydomain.com
      http:
        paths:
          - path: /bar/*
            backend:
              serviceName: bar
              servicePort: 8080
```

# Kubernetes - Setup with kubeadm

- 3 Ubuntu 16.04 servers with 4GB RAM and private networking enabled
- On each of the three Ubuntu 16.04 servers run the following commands as root:

```
apt-get update && apt-get install -y apt-transport-https
```

```
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -
```

```
cat <<EOF >/etc/apt/sources.list.d/kubernetes.list
```

```
deb http://apt.kubernetes.io/ kubernetes-xenial main
```

```
EOF
```

```
apt-get update
```

```
apt-get install -y kubelet=1.18.0-00 kubeadm=1.18.0-00 kubectl=1.18.0-00 docker.io
```

# Kubernetes - Setup with kubeadm

## Setup the Kubernetes Master

- `kubeadm init`
- `mkdir -p $HOME/.kube`
- `sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config`
- `sudo chown $(id -u):$(id -g) $HOME/.kube/config`

## Join your nodes to your Kubernetes cluster

- `kubeadm join --token 702ff6.bc7aacff7aacab17 174.138.15.158:6443  
--discovery-token-ca-cert-hash  
sha256:68bc22d2c63180ofd358a6d7e3998e598deb2980ee613b3c2f1da8978960c8ab`

## Installing the Calico Add-On

- `curl https://docs.projectcalico.org/manifests/calico.yaml -O`
- `kubectl apply -f calico.yaml`